



OpenCL for realtime graphics

Matthäus G. Chajdas



Why OpenCL?

- Because OpenCL does it all
 - **Single source** for **all** graphics APIs (D3D9, D3D10, D3D11, OpenGL 2.x – 4.x)
 - Mobile support
 - Web support
- Clean interop API
- Easy debugging (printf on CPU!)

Interop API

- Acquire resource, map/unmap
- Automatic synchronization
- Minimal code required
 - Full production implementation: 900 LoC
 - Includes error checking & comments
 - Supports OpenCL 1.1 & 1.2, OpenGL, Direct3D11 KHR/NV

Tile based deferred shading

Tile-based Deferred Shading

1. Divide screen into tiles and determine which lights affects which tiles

2. Only apply the visible light sources on pixels

- › Custom shader with multiple lights
- › Reduced bandwidth & setup cost

How can we do this best in DX11?



Tile based deferred shading

Tile-based Deferred Shading

1. Divide screen into tiles and determine which lights affects which tiles


2. Only apply the visible light sources on pixels

- › Custom shader with multiple lights
- › Reduced bandwidth & setup cost

How can we do this best in



Tile based deferred shading




Tile based deferred shading

- Works!
- Performance roughly the same as DirectCompute
- Some limitations apply though ...

Tile based deferred shading

- No image read/write
 - Just map image twice & get the desired undefined behavior :)
 - Fixed in OpenCL 2.0
- No depth images
 - Fixed by extension (for OpenGL)
- No MSAA images
 - Fixed by extension (for OpenGL)

230.84 ms (4.3 fps)



Hawaii

Real compute

- Voxel raytracing
 - 600 LoC of kernel code
 - Tree traversal
 - Per-thread stacks
- Run out of GPU memory?
 - Current GPUs max out at 16 GiB (AMD) or 12 GiB (NVIDIA)
 - CPUs max out at ... well ...

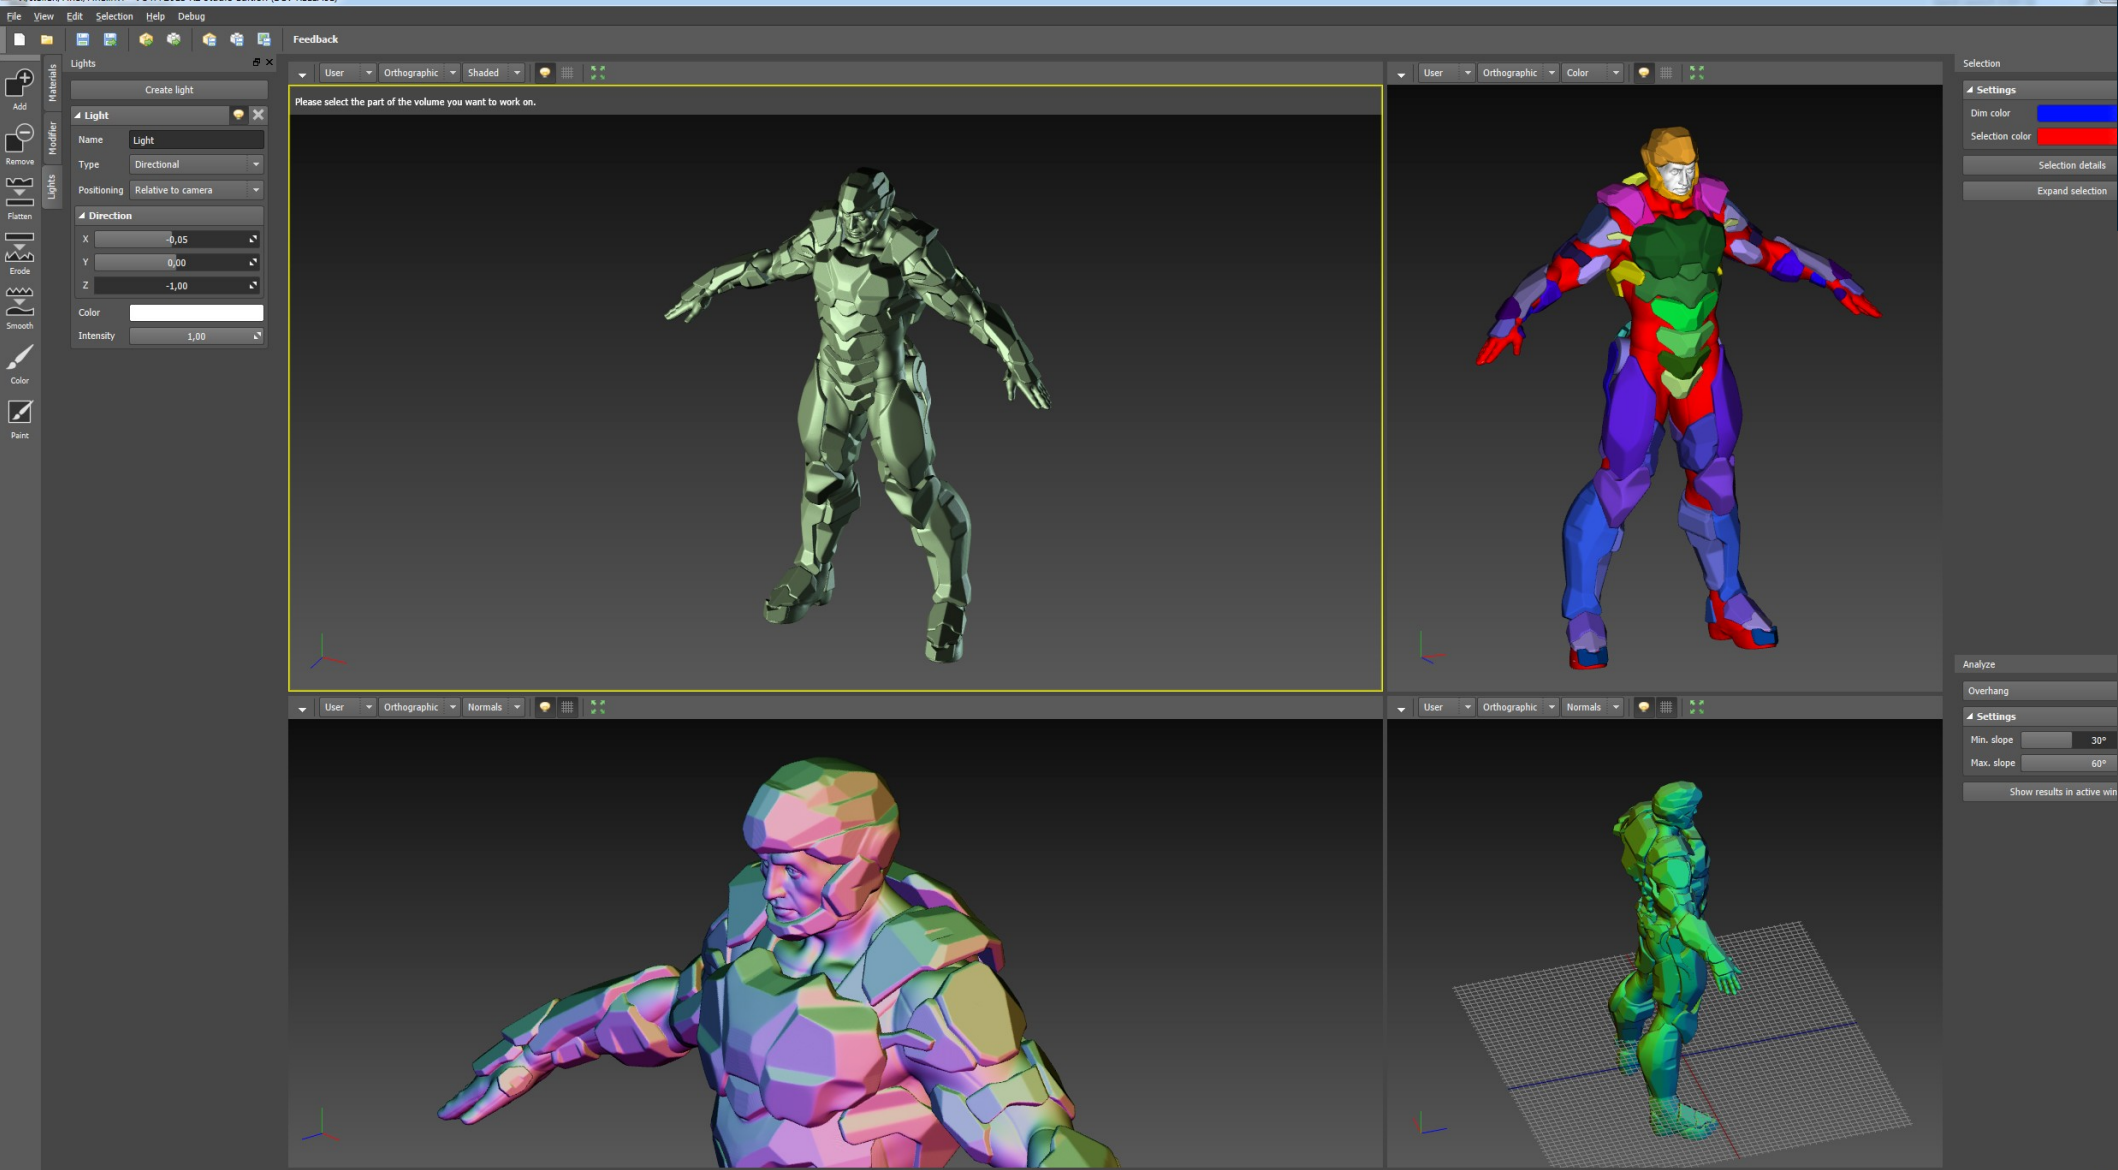
Voxel raytracing

- Same code
- Really!
- Portable performance is not a myth!



Voxel raytracing

- GPU optimizations benefit CPU and vice versa
- No special casing
 - Avoid some “obviously” slow paths
 - AMD & NVIDIA* get comparable performance/efficiency
 - CPU performance ~ 30% slower than extremely optimized ISPC code



Shipping products/VOTA

- Voxel editor
- All editing is done in OpenCL only
- Graphics interop running all the time
 - OpenCL prepares data for visualization
 - Minimal CPU intervention (to minimize buffer sizes)

VOTA, lessons learned

- All users have OpenCL installed
- AMD doesn't have a kernel cache (just do it on your own)
- Compiler quality is “good enough”
 - Still, we eagerly wait for SPIR
 - Initial startup takes quite some time due to compiling
- Query everything which is implementation specific

OpenCL/Graphics interop state

- AMD
 - Direct3D: Works
 - OpenGL: Works mostly* (Linux & Windows)
- NVIDIA
 - Direct3D: Through _NV extension, _KHR not supported.
 - OpenGL: Works (Linux & Windows)

OpenCL/Graphics interop state

- AMD's implementation is much stricter (similar to OpenGL)
- Both have occasional issues with memory management
 - Need extension to force GPU memory defragment & query where buffer ended up
- For debugging/profiling, use CodeXL/KernelAnalyzer
 - NVIDIA stopped supporting OpenCL profiling some time ago

Reporting bugs

- AMD: Just report through normal channels. They got back to me quickly & and you get status updates (cool)
- NVIDIA, best report through CUDA bug report form. They don't say much about it, but they still fix OpenCL bugs.
- Intel: The forums seem to work fine.

OpenCL for graphics, the road ahead

- With OpenCL 2.0, OpenCL will far surpass DirectCompute/OpenGL compute shaders
 - GPU pipelining
 - Shared CPU/GPU memory data structures
 - AMD is aggressively working on it, first OpenCL 2.0 SDK should come this June
 - Intel most likely working on it as well

Conclusion

- You can ship an application built on OpenCL **today**
- For games, you'll likely want the MSAA/depth extensions
 - Vendors should do it instantly **once requested**
 - Can start writing the game today without problems
- In future, hard to avoid
 - OpenCL 2.0 allows GPU to traverse scene graph, prepare draw commands, then issue draw calls from CPU, with zero copy overhead

Thanks to ...

- The whole Volumerics team
 - @janjorrit, @mreitinger
- @repi
- AMD
 - @grahamsellers, @JCBaratault
- Codeplay & Khronos
 - @codeandrew, @neilt3d